



Azure Sentinel Level 400 KQL Workshop



[Aka.ms/sentinelPerth](https://aka.ms/sentinelPerth)

New Query 1*

HelpSettingsSample queriesQuery explore

CyberSecurityDemo

Run

Time range: Last 24 hours

SaveCopyExportNew alert rulePin to dashboard

SchemaFilter

Filter by name or type...

Collapse all

Active

CyberSecurityDemo

DnsAnalyticsLogManagementOffice365Security

CommonSecurityLogLinuxAuditLogProtectionStatusSecurityAlert

AlertNameAlertSeverityConfidenceLevelConfidenceScoreDescriptionDisplayNameEndTimeEntitiesExtendedLinksExtendedPropertiesIsIncident

SecurityAlert

limit 50

Queries

Completed. Showing results from the last 24 hours.

00:00:03.05543 recordsDisplay time (UTC+00:00)

TABLECHARTColumns

Drag a column header and drop it here to group by that column

	TimeGenerated [UTC]	DisplayName	AlertName
2019-07-08T18:56:26.000	Anonymous IP address	Anonymous IP address	
...			
TenantId	ab86c959-1ba3-495c-a00d-ced30d8825d3		
TimeGenerated [UTC]	2019-07-08T18:56:26Z		
DisplayName	Anonymous IP address		
AlertName	Anonymous IP address		
AlertSeverity	Medium		
Description	Sign-in from an anonymous IP address (e.g. Tor browser, anonymizer VPNs)		
ProviderName	IPC		
VendorName	Microsoft		
VendorOriginalId	7fb7db35c22be5f914f827daa5d29d39c65de080d554781eb3221014f8690f9a		
SystemAlertId	e6692353-f03c-489a-a0fc-14cae4bae0e2		
AlertType	AnonymousLogin		

1 of 150 items per page1 - 43 of 43 items

KQL Column Types

Basic

- **int, long**
- **bool**: true, false
- **string**: "example", 'example'

Time

- **datetime**: datetime(2016-11-20 22:30:15.4), now(), ago(4d)
- **timespan**: 2d, 20m, time(1.13:20:05.10), 100ms

Complex

- **dynamic**: JSON format

'where' command

Filters a table to the subset of rows that satisfy a predicate.

Syntax: *T* | *where* *Predicate*

Examples: *SecurityEvent* | *where* *TimeGenerated* > *ago(1d)*

SecurityEvent | *where* * *contains* "Kusto"

- String predicates: ==, has, contains, startswith, endswith, matches regex, etc
- Numeric/Date predicates: ==, !=, <, >, <=, >=
- Empty predicates: isempty(), notempty(), isnull(), notnull()

'where' exercise

```
SecurityEvent  
| where TimeGenerated > ago(1d)
```

```
SecurityEvent  
| where TimeGenerated > ago(1h) and EventID == 4624 // Successful logon
```

```
SecurityEvent  
| where TimeGenerated > ago(1h)  
| where EventID == 4624  
| where AccountType =~ "user" // case insensitive
```

```
Perf  
| where InstanceName matches regex "^[A-Z]:"
```

'limit' / 'take' command

Return up to the specified number of rows.

Syntax: *T* | *limit* <number>

Example: *SecurityEvent* | *limit* 5

- Sort is not guaranteed to be preserved.
- Consistent result is not guaranteed (when running the same query twice)
- Very useful when trying out new queries.
- Default limit is 10,000.

'limit' exercise

SecurityEvent

```
| limit 10
```

SecurityEvent

```
| where TimeGenerated > ago(1h)
```

```
| where EventID == 4624
```

```
| where AccountType =~ "user"
```

```
| take 10
```


'count' command

Returns the number of records in the input record set.

Syntax: *T* | *count*

Example: *SecurityEvent* | *count*

'count' exercise

Perf

```
| count
```

SecurityEvent

```
| where TimeGenerated > ago(1h)
```

```
| where EventID == 4624
```

```
| count
```

'summarize' command

Produces a table that aggregates the content of the input table.

Syntax: *T | **summarize** Aggregation [by Group Expression]*

Examples: *SecurityEvent | **summarize** count() by Computer*

- Simple aggregation functions: count(), sum(), avg(), min(), max(),
- Advanced functions (next slide): arg_min(), arg_max(), percentiles(), makelist(), countif()
- No Group Expression implies 'distinct'.

'summarize' exercise

Perf

```
| where CounterName == "Free Megabytes"  
| where InstanceName matches regex "^[A-Z]:$"  
| summarize min(CounterValue)
```

SecurityEvent

```
| where TimeGenerated > ago(1h)  
| where EventID == 4624  
| summarize count() by AccountType, Computer
```

'summarize' advanced aggregations

arg_min(), arg_max(): returns the extreme value

Example: *Supplies | **summarize arg_min**(Price, Supplier) by Product*
// Cheapest supplier of each product

percentiles(): returns the value at the percentile

Example: *CallDetailRecords | **summarize percentile**(Duration, 95) by continent*
// The value of Duration that is larger than 95% of the sample set

makelist(), makeset(): returns a list of all values/distinct values respectively

Example: *PageViewLog | **summarize** countries=**make_set**(country) by continent*

'summarize' advanced aggregations exercise

SecurityEvent

```
| where EventID == 4624  
| summarize arg_max(TimeGenerated, *) by Account
```

AzureDiagnostics

```
| summarize arg_max(TimeGenerated, *) by ResourceId
```

SecurityEvent | summarize

AdminSuccessfulLogons =

countif(Account contains "Admin" and EventID == 4624),

AdminFailedLogons =

countif(Account contains "Admin" and EventID == 4625)

Quiz #1

What is the difference between the following queries?

SecurityEvent

```
| summarize arg_max(TimeGenerated, *) by Account  
| where EventID == "4624"  
| count
```

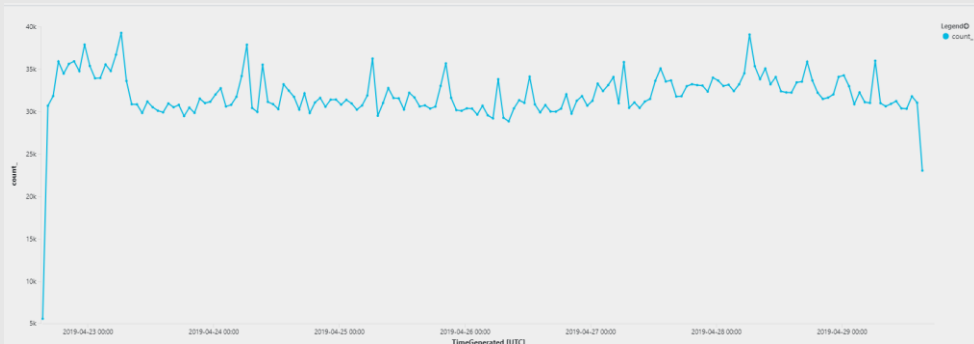
SecurityEvent

```
| where EventID == "4624"  
| summarize arg_max(TimeGenerated, *) by Account  
| count
```

'summarize': bin and time series

A very useful summarize operation is creating time series:

`SecurityEvent | summarize count() by bin(TimeGenerated, 1h) | render timechart`



Other time measurements: 1h, 5d, 10m (defaults to 1h)

Can create multiple legends by aggregating additional field

'bin' exercise

SecurityEvent

```
| where TimeGenerated > ago(7d)  
| summarize count() by bin(TimeGenerated, 1d)
```

Perf

```
| where CounterName == "Free Megabytes"  
| where InstanceName matches regex "[A-Z]:$"  
| summarize min(CounterValue) by bin(TimeGenerated, 1d)
```

Question #1

Use the [open to use KQL playground](https://aka.ms/LAdemo) (<https://aka.ms/LAdemo>)

1. Find how many times each process ran per computer (hint: look for event ID that represent new process creation)
2. Render a time chart of free MB over last 7 days for disk C: (hint: InstanceName should be 'C:', CounterName should be 'Free Megabytes')

#1 solution

//Question 1

```
SecurityEvent  
| where EventID == 4688  
| summarize count() by NewProcessName, Computer
```

//Question 2

```
Perf  
| where InstanceName == "C:"  
| where CounterName == "Free Megabytes"  
| summarize FreeMB = avg(CounterValue) by bin(TimeGenerated, 1h) | render  
timechart
```

'extend' command

Create calculated columns and append them to the result set.

Syntax: *T | extend ColumnName [= Expression] [, ...]*

Example: *SecurityEvent | extend ComputerNameLength = strlen(Computer)*

- The new added column is not indexed.
- To only change a column name, use 'project-rename'.
- Useful function for in 'extend': iff, extract

'project' command

Select the columns to include, rename or drop, and insert new computed columns.

Syntax: *T* | *project* *ColumnName* [= *Expression*] [, ...]

Example: *SecurityEvent* | *project* *TimeGenerated*, *Computer*

'| project-away' – Removed specified column/s.

'| project-rename' – Rename specified column/s.

'extend' & 'project' exercise

Perf

```
| where CounterName == "Free Megabytes"  
| where InstanceName == "C:"  
| extend FreeKB = CounterValue * 1000  
| extend FreeGB = CounterValue / 1000
```

Perf

```
| where CounterName == "Free Megabytes"  
| project Computer , CounterName , CounterValue
```

Perf

```
| where CounterName == "Free Megabytes"  
| extend FreeKB = CounterValue * 1000  
| extend FreeGB = CounterValue / 1000  
| extend FreeMB = CounterValue  
| project Computer , CounterName , FreeGB , FreeMB , FreeKB
```

'distinct' command

Produces a table with the distinct combination of the provided columns of the input table.

Syntax: *T* | *distinct* *Column1*, *Column2*

Example: *SecurityEvent* | *distinct* *Computer*

'order by' / 'sort by' & 'top' operator

Order by: Sort the rows of the input table into order by one or more columns.

Top: returns the top values after sort. Faster and can sort by expression

Syntax: *T | sort by column [asc | desc] [nulls first | nulls last]*

T | top NumberOfRows by Expression [asc | desc] [nulls first | nulls last]

Example: *Table | order by country asc, price desc*

Don't assume order by default

'order by' / 'top' exercise

SecurityEvent

```
| where TimeGenerated > ago(7d)  
| order by TimeGenerated desc  
| limit 100 // try also top
```

SecurityEvent

```
| top 100 by TimeGenerated desc
```

SecurityEvent

```
| where EventID == 4624  
| summarize cnt=count() by Account  
| top 10 by cnt
```

Question #2

Use the [open to use KQL playground](https://aka.ms/LAdemo) (<https://aka.ms/LAdemo>)

The table "SecurityEvent" contains security events collected from Windows machines. The events are identified by their IDs. Find the column that represents this ID, and complete the following tasks:

- Render graph of logon events over time, starting from two weeks ago until one week ago.
- Render graph of successful (4624) vs failed (4625) logons over the last 7 days, use alias for the legend ("Success", "Failed")
- Find the last logon time for a user named "ContosoASCAAlert\\LBecker"

#2

//Question #A: Render graph of logon events over time, starting from two weeks ago until one week ago.

SecurityEvent

```
| where EventID == 4624
| summarize count() by bin(TimeGenerated, 1h)
| render timechart
```

//Question #B: Render graph of successful (4624) vs failed (4625) logons over the last 7 days, use alias for the legend ("Success", "Failed")

SecurityEvent

```
| summarize Success=countif(EventID == 4624), Failed=countif(EventID==4625) by
bin(TimeGenerated, 1h)
| render timechart
```

//Question #C: Find the last logon time for a user named "ContosoASCAAlert\\LBecker"

SecurityEvent

```
| where EventID == 4624
| where Account == "ContosoASCAAlert\\LBecker"
| summarize max(TimeGenerated)
```

'extract' function

Get a match for a regular expression from a text string.

Syntax: *extract(regex, captureGroup, text [, typeLiteral])*

Example: *extract("x=([0-9.]+)", 1, "hello x=45.6|wo") == "45.6"*

'let' statement

Let statements bind names to expressions.

- declare global 'variable' or reuse of 'variables'.
- Reuse of query runs
- Declare functions
- Declate dynamic table

Example: See on demo.

'union' operator

Takes two or more tables and returns the rows of all of them.

Example:

SecurityEvent | union (SecurityAlert | where Severity > 3)

- kind=inner(common columns), outer (all columns- default)
- Supports wildcard to union multiple tables (union Security*)
- Can union between tables from different clusters (or workspaces)

'union' exercise

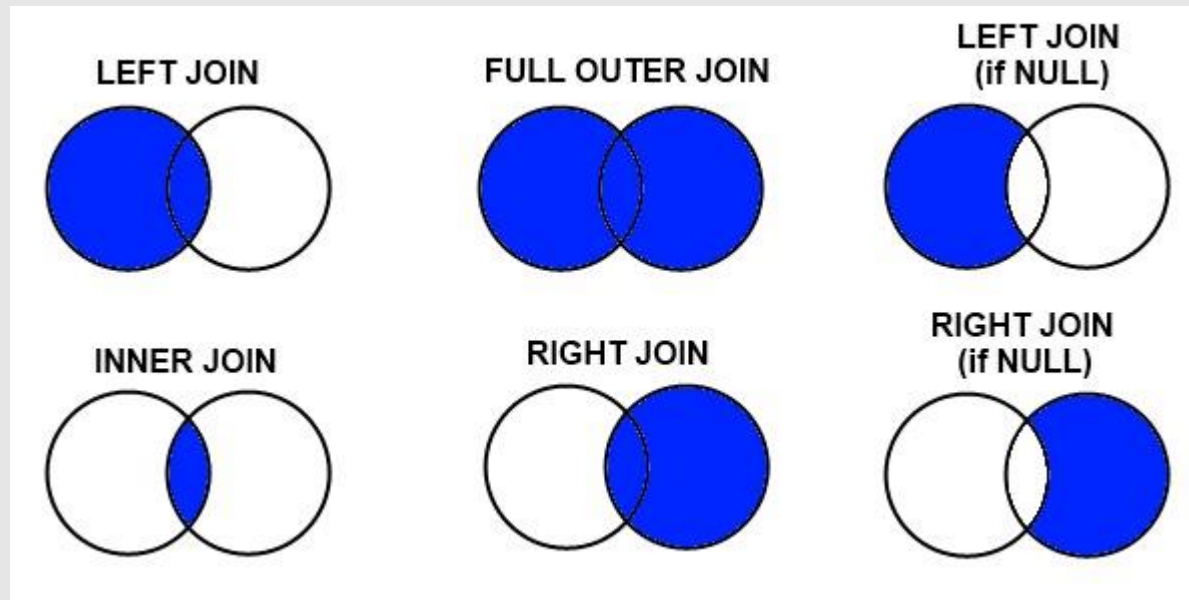
```
SecurityEvent  
| union Heartbeat  
| summarize count() by Computer
```

'join' operator

Merge the rows of two tables to form a new table by matching values of the specified column(s) from each table.

Syntax: *LeftTable* | join [*JoinParameters*] (*RightTable*) on *Attributes*

Example: *SecurityEvent* | join (*SecurityAlert* | *where Severity > 3*) on *Account*



Question #3

Find the ratio of alerts (in the SecurityAlert table) to events (in the SecurityEvent table)

#3 solution

```
SecurityAlert
| extend table = "SecurityAlerts"
| union (SecurityEvent | extend table = "SecurityEvents")
| summarize SecurityAlerts = countif(table == "SecurityAlerts") ,
SecurityEvents = countif(table == "SecurityEvents")
| extend Ratio = SecurityAlerts * 1.0 / SecurityEvents
| project SecurityEvents , SecurityAlerts , Ratio
```